

Nimble Cheatsheet

Nimble is built on top of Python's most popular data science and machine learning libraries to provide a single, easy to use, API for any data science job.

Nimble Data Object

Nimble has 4 data types that share the same API.

Each use a different backend to optimize the operations based on the type of data in the object. Choosing the type that best matches the data will support more efficient operations.

<u>Type</u>	<u>Data</u>	<u>Backend</u>
List	any data	Python list
Matrix	all the same type	NumPy array
DataFrame	each column has 1 type	Pandas DataFrame
Sparse	mostly missing or 0	SciPy coo_matrix

Visualization of a Nimble data object:

```
graph TD; features --> size; features --> span_speed_class[span speed class]
```

The diagram illustrates the structure of a Nimble data object. It shows a box labeled "features" at the top, which has two arrows pointing down to two separate boxes. The left box is labeled "size" and the right box is labeled "span speed class".

points

features

I/O

Creating Data

`nimble.data` is the primary function for loading data from all accepted sources. It accepts raw python objects, strings that are paths to files or urls, and open file objects.

```
X = nimble.data('DataFrame', [[1, 'a'], [2, 'b']])
```

From convenience, `nimble.ones`, `nimble.zeros`, and `nimble.identity` are available to quickly create objects with specific data. The following create objects with 10 points and 10 features.

```
allOnes = nimble.ones('Matrix', 10, 10)
allZeros = nimble.zeros('Sparse', 10, 10)
identity = nimble.identity('List', 10)
```

`nimble.random.data` is available to construct an object of random data with adjustable sparsity. The following creates a Matrix object with 10 points, 10 features and 0 sparsity.

```
randomData = nimble.random.data('Matrix', 10, 10, 0)
```

Fetching

Fetching returns the local path(s) to an online dataset, downloading and saving it necessary.

```
fileLocation = nimble.fetchFile('https://link.to.dataset.csv')
fileLocationsList = nimble.fetchFiles('UCI::iris')
```

Saving

Nimble data objects can be written to a csv or mtx file or saved as a pickle file. `TrainedLearner` objects can also be pickled.

```
X.writeFile('saved.csv')
X.save('saved.nimbd')
trainedLearner.save('saved.nimmm')
```

Information about the data

Some information is set automatically on creation. By default automatic detection of pointNames and featureNames occurs. Data information can also be controlled by some of the parameters for [nimble.data](#).

```
>>> X = nimble.data('DataFrame', '/path/to/X.csv')
>>> X.shape           # always set
(3, 4)
>>> X.path            # set when source is a path
'/path/to/X.csv'
>>> X.features.getNames() # automatically detected
['h', 'w', 'a']
>>> X.points.getNames()  # automatically detected
['0k1r3', '6t3n1', '8i7g3', '0k2r2']
>>> headers = ['height', 'width', 'depth']
>>> items = ['couch', 'table', 'chair', 'love seat']
>>> X = nimble.data('Matrix', '/path/to/dataset.csv',
...               pointNames=items, featureNames=headers,
...               name='furniture')
```

Once the object is created, the object's methods can be used to get or set information about the object.

```
X.name           # getter and setter
X.absolutePath  # getter only
X.relativePath  # getter only
X.[points/features].getNames()
X.[points/features].getName(index)
X.[points/features].setNames(assignments)
X.[points/features].setName(oldIdentifier, newName)
X.[points/features].getIndex(identifier)
X.[points/features].getIndices(names)
X.[points/features].hasName(name)
```

The methods of the Nimble data object control operations that apply to the entire object or the elements. The points and features properties of the object have additional methods for operations that apply along that axis of the data object.

Note: Nimble can also be used when your data points are each matrices or higher dimensional objects.

Note: Nimble can also be used when your data points are each matrices or higher dimensional objects.

Visualization	
Printing	
Nimble provides several ways to print or stringify the data, with varying levels of flexibility.	
# a representation of the data object that conforms to Python's repr standards # a pretty-printed representation of the data object # pretty-print the object with customized parameters	
Plotting	
Nimble provides basic plotting functions using the matplotlib package on the backend.	
# A scatter plot showing one feature plotted against another feature # A rolling average of one feature plotted against another feature # Plot a histogram of the distribution of values in a feature # Plot the means of a feature grouped by another feature # Plot an aggregate statistic for each group of a feature # Display a heat map of the data # Bar chart comparing points/features # Plot means with 95% confidence interval bars # Bar chart comparing an aggregate statistic between points/features	
Iteration	
Iteration can occur over elements, points, or features.	
# a single value # new Nimble data object containing the data from a single point # new Nimble data object containing the data from a single feature	

```
X
print(X)
X.show(description, ...)
>>> for element X.iterateElements(order, only):
...     print(element)
>>> for point in X.points:
...     print(point)
>>> for feature in X.features:
...     print(feature)

X.plotFeatureAgainstFeature(x, y, ...)
X.plotFeatureAgainstFeatureRollingAverage(x, y, ...)
X.plotFeatureDistribution(feature, ...)
X.plotFeatureGroupMeans(feature, groupFeature, ...)
X.plotFeatureGroupStatistics(statistic, feature,
                             groupFeature, ...)

X.plotHeatMap(...)
X[points/features].plot(identifiers, ...)
X[points/features].plotMeans(identifiers, ...)
X[points/features].plotStatistics(statistic,
                                  identifiers,
                                  ...)
```

Querying	
Data Querying	
Many methods provide information about the data within a Nimble data object. The following functions provide information or perform calculations on the data, but they do not modify the data in the object or return a new Nimble data object.	
# the number of elements satisfying the query # values and counts of unique elements # True if any elements are equal to zero, otherwise False # information describing the contents of the object # number of points/features satisfying the query # identify points/features satisfying the query # similarity calculations on each point/feature # statistics calculations on each point/feature # removal of duplicate points/features # statistical information about each feature	

```
X.countElements(condition)
X.countUniqueElements(...)
X.containsZero()
X.report()
X[points/features].count(condition)
X[points/features].matching(function)
X[points/features].similarities(function)
X[points/features].statistics(function, ...)
X[points/features].unique()
X.features.report(basicStatistics,
                  extraStatisticFunctions)
```

Query Strings
A string Nimble uses to create a function from comparison operators (==, !=, >, <, >=, <=) or "is" or "is not" and a nimble.match function or Python's True, False, or None. See the QueryString object.
Element Query numGreaterThan10 = X.countElements("> 10") numNonMissing = X.countElements("is not missing")
Axis Query (using feature names from the example) bigSpan = X.points.count("span > 30") eagles = X.points.extract("class == eagle") fast = X.points.copy("speed > 200")

Indexing
Nimble uses INCLUSIVE indexes to support consistent behavior when using names or indices as identifiers. Indexing can be performed from the data object or the points and features attributes.
data['bird2', 'speed'] data[1, 2] data['bird2':'bird4', [0, 2]] X.features["span"] X.features[2] X.points['bird4'] X.points[3] X.features[:'speed'] X.points[3:]

Data Manipulation

↘ indicates an in-place operation that modifies the original data object rather than returning a copy

Math

~~Python operators~~ be used between a Nimble data object and either ~~or two~~ Nimble data objects. The objects must be the same shape for elementwise operations and compatible shapes for matrix multiplication.

X + Y	Elementwise Addition	X ** Y	Elementwise Power
X - Y	Elementwise Subtraction	X % Y	Elementwise Modulo
X * Y	Elementwise Multiplication	X @ Y	Matrix Multiplication
X / Y	Elementwise Division		

The `stretch` property allows for expanded (broadcasting) computation with one-dimensional data objects. The one-dimensional object is stretched (repeated) to match the shape of the other object.

X + Y.stretch	# 2D + 1D
X.stretch / Y	# 1D / 2D
X.stretch * Y.stretch	# 1D * 1D

Linear algebra functions can also be applied to Nimble data objects.

X.matrixMultiply(other)	# (same as using @ operator)
X.matrixPower(power)	# a square matrix times itself 'power' times
X.inverse()	# the inverse of the matrix
X.solveLinearSystem(b)	# find the solution to a linear system
X.T	# returns the transposed object

Copying and Reordering

X.copy(to)	
X[points/features].copy(toCopy, ...)	
↘ X[points/features].permute(order)	
↘ X[points/features].sort(by, ...)	

Data Cleaning

Element Modification

↘ X.replaceFeatureWithBinaryFeatures(featureToReplace)	# replace a categorical feature with one-hot encoded features
↘ X.replaceRectangle(replaceWith, pointStart, featureStart, ...)	# replace a section of the data with other data
↘ X.transformElements(toTransform, ...)	# change elements to new values
X.calculateOnElements(toCalculate, ...)	# apply a calculation to each element
↘ X.transformFeatureToIntegers(featureToConvert)	# map unique values to an integer and replace each element with the integer value
↘ X[points/features].fillMatching(fillWith, matchingElements, ...)	# replace elements in points/features with a different value(s)
↘ X[points/features].replace(data, ...)	# replace points/features with a new points/features
↘ X[points/features].normalize(function, ...)	# replace elements in points/features with the normalized equivalents
↘ X[points/features].transform(function, ...)	# modify the elements within points/features
X[points/features].calculate(function, ...)	# apply a calculation to the elements within points/features

Structural Changes

↘ X.transpose()	# invert the points and features of this object (inplace)
↘ X.flatten(order, ...)	# deconstruct this data into a single point
↘ X.unflatten(dataDimensions, order, ...)	# expand a one-dimensional object into a new shape
X.groupByFeature(by, ...)	# separate the data into groups based on the value in a single feature
↘ X[points/features].append(toAppend)	# add additional points/features to the end of the object
↘ X[points/features].insert(insertBefore, toInsert, ...)	# add additional points/features at a given index
↘ X[points/features].extract(toExtract, ...)	# remove points/features from the object and place them in a new object
↘ X[points/features].delete(toDelete, ...)	# remove points/features from the object
↘ X[points/features].retain(toRetain, ...)	# keep certain points/features of the object
X[points/features].mapReduce mapper, reducer)	# apply a mapper and reducer function to each point/feature
X[points/features].repeat(totalCopies, copyOneByOne)	# make a repeated copies of the object

Training, Applying, and Testing	Interfaces
<p>The same API is available for any available learner.</p> <pre>trainedLearner = nimble.train(learnerName, trainX, trainY, ...) # Learn from the training data # Returns a TrainedLearner predictedY = nimble.trainAndApply(learnerName, trainX, trainY, testX, ...) # Make predictions on new data performance = nimble.trainAndTest(trainX, trainY, testX, testY, # Evaluate the accuracy of the predictions on the performanceFunction, ...) # testing data performance = nimble.trainAndTestOnTrainingData(trainX, trainY, # Evaluate the accuracy of the predictions on the performanceFunction, ...) # data used for training kFoldCrossvalidator = nimble.crossValidate(learnerName, X, Y, # Evaluate the accuracy with varying arguments performanceFunction, ...) # Returns a KFoldCrossvalidator normalizedX = nimble.normalizeData(learnerName, trainX, ...) # Transform the data through normalization filledX = nimble.fillMatching(learnerName, matchingElements, trainX, ...) # replace matching elements in points/features with # provided or calculated values</pre>	<p>Nimble interfaces with popular machine learning packages, to apply their algorithms within our API. Interfaces are used by providing "package.learnerName". For example:</p> <pre>nimble.train("nimble.RidgeRegression", ...) nimble.trainAndApply("sklearn.KNeighborsClassifier", ...) nimble.trainAndTest("keras.Sequential", ...)</pre> <p>Nimble interfaces with popular machine learning packages, to apply their algorithms within our API. Interfaces are used by providing "package.learnerName". For example:</p> <pre>nimble.showAvailablePackages() nimble.learnerNames() nimble.showLearnerNames()</pre>
Learner Arguments	TrainedLearner
<p>To find the parameters and any default values for a learner.</p> <pre>nimble.learnerParameters(name) # A list of parameters that the learner accepts nimble.showLearnerParameters(name) # Print parameters of the learner nimble.learnerParameterDefaults(name) # A dictionary of parameters and their default values nimble.showLearnerParameterDefaults(name) # Print the default values of the learner</pre> <p>Arguments can be set in two ways. Using the arguments parameter in the nimble function or passing the learner object's parameters as keyword arguments.</p> <pre>>>> t1 = nimble.train("sklearn.KNeighborsClassifier", trainX, trainY, arguments={'n_neighbors': 7}) >>> t1 = nimble.train("sklearn.KMeans", trainX, trainY, n_clusters=7)</pre> <p>Cross-validation can be triggered using a nimble.CV object.</p> <pre>>>> t1 = nimble.train("sklearn.Ridge", trainX, trainY, alpha=nimble.CV([0.1, 1.0]))</pre> <p>When a package requires another object from the package, use nimble.Init with the arguments required to instantiate the object.</p> <pre>>>> layer0 = nimble.Init('Dense', units=64, activation='relu', input_dim=256) >>> layer1 = nimble.Init('Dropout', rate=0.5) >>> layer2 = nimble.Init('Dense', units=10, activation='softmax') >>> t1 = nimble.train("keras.Sequential", trainX, trainY, layers=[layer0, layer1, layer2], ...)</pre>	<p>The nimble.train function returns a TrainedLearner (referred to as "t1" below).</p> <pre>t1.learnerName # The name of learner used for training t1.arguments # The arguments used for training t1.randomSeed # The randomSeed applied for training t1.crossValidation # KFoldCrossvalidator object of the cross-validation results t1.apply(testX, ...) # Apply the trained learner to new data t1.getAttributes() # Dictionary with attributes generated by the learner t1.getScores(testX, ...) # The scores for all labels for each data point t1.incrementalTrain(trainX, trainY, ...) # Continue to train with additional data t1.retrain(trainX, trainY, ...) # Train the learner again on different data t1.save(outPath) # Save the learner for future use. t1.test(testX, testY, # Evaluate the accuracy of the learner on testing data performanceFunction, ...)</pre>
Helper Modules	
<pre>nimble.calculate # Common calculation functions such as statistic and performance functions nimble.match # Common functions for determining if data satisfies a certain condition nimble.fill # Common functions for replacing missing data with another value nimble.random # Support for random data and random control within Nimble nimble.learners # Nimble's prebuilt custom learner algorithms nimble.exceptions # Nimble's custom exceptions types</pre>	